

# Data Science/Practical ML exercise

*Nikolay Dobrinov*

*Oct 18, 2017*

## Executive Summary

“One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. More information is available from the website here: <http://web.archive.org/web/20161224072740/http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset). The training and testing data is located here:

- <https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>
- [https://d396qusza40orc.cloudfront.net/pml-testing.csv](https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv) "

The data was pre-processed and four classification methods were tested. Random forests performs best, with 99.54% accuracy on 100 trees. The most important 5 variables in the random forest model are roll.belt, yaw.belt, magnet.dumbbell.z, magnet.dumbbell.y, pitch.belt.

## Data pre-processing

The seed is fixed for the purposes of reproducibility. We split the training data into a train and test data sets to cross validate the models. The testing data set provided for this assignment is used as a validation set.

```
setwd("/Users/nikolaydobrinov/Documents/work/Courses/R/WorkDirectory/Course8_week4_coding_assignments")

library(dplyr)
library(caret)

set.seed(333) # use a seed for replicability

# load data
traintest <- read.csv("./data/pml-training.csv", na.strings = c("NA", ""))
validate <- read.csv("./data/pml-testing.csv", na.strings = c("NA", ""))

# split train-test
inTrain <- createDataPartition(y=traintest$classe, p=0.7, list=FALSE)
train <- traintest[inTrain,]
test <- traintest[~inTrain,]
```

Remove variables that do not seem useful, or it is not clear what they represent. Note that user\_name should not be used in classification as the prediction algorithm should work regardless of the specific user using the device Remove the first 7 columns

```
train <- select(train, -(1:7))
test <- select(test, -(1:7))
validate <- select(validate, -(1:7))
```

Remove variables with NAs. naVars below reveals that in all variables where NAs exist, about 98% of the observations are NA. We remove all of these features

```
naVars <- sapply(train, function(x) sum(is.na(x)))/nrow(train)
naVarsExclude <- names(naVars[naVars > 0])
train <- train[, !names(train) %in% naVarsExclude]
test <- test[, !names(test) %in% naVarsExclude]
validate <- validate[, !names(validate) %in% naVarsExclude]
```

Check for variables with low variation and remove them. There are no variables with low or zero variance

```
lowVariance <- nearZeroVar(train, saveMetrics=TRUE)
sum(lowVariance$zeroVar) + sum(lowVariance$nzv)
```

```
## [1] 0
```

## Analysis

I fit four classification methods - tree, random forest, generalized boosted regression (gbm), and linear discriminant analysis (lda). The corresponding classification accuracy of the models on the testing sample is: single tree - 49%; random forest with 100 trees - 99.34%; gbm - 96%; lda - 71%. This section presents the results on the lowest error model, random forests, the rest of the models are presented in the Appendix.

The random forest model with 100 trees produces an expected error on the test sample of 0.66%. The variable importance function reveals that across the 100 trees the most important 5 variables are related to the belt and the position of the dumbbell: roll.belt, yaw.belt, magnet.dumbbell.z, magnet.dumbbell.y, pitch.belt. Optimal number of variables to be randomly sampled as candidates at each split is mtry=2.

```
pred.rf <- predict(modFit.rf, test) # predict on test data
confusionMatrix(pred.rf, test$classe) # measure accuracy on test data
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1673    3    0    0    0
##           B    0 1135    7    0    0
##           C    1    1 1017   19    0
##           D    0    0    2  944    5
##           E    0    0    0    1 1077
##
## Overall Statistics
##
##           Accuracy : 0.9934
##           95% CI : (0.991, 0.9953)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9916
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
```

```
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9994  0.9965  0.9912  0.9793  0.9954
## Specificity      0.9993  0.9985  0.9957  0.9986  0.9998
## Pos Pred Value   0.9982  0.9939  0.9798  0.9926  0.9991
## Neg Pred Value   0.9998  0.9992  0.9981  0.9959  0.9990
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2843  0.1929  0.1728  0.1604  0.1830
## Detection Prevalence 0.2848  0.1941  0.1764  0.1616  0.1832
## Balanced Accuracy 0.9993  0.9975  0.9935  0.9889  0.9976
```

```
varImp(modFit.rf) # variable importance
```

```
## rf variable importance
##
## only 20 most important variables shown (out of 52)
##
##          Overall
## roll_belt      100.00
## yaw_belt       73.83
## magnet_dumbbell_z 65.35
## pitch_forearm  59.97
## magnet_dumbbell_y 58.36
## pitch_belt     55.70
## magnet_dumbbell_x 51.32
## roll_forearm   48.04
## accel_belt_z   43.22
## roll_dumbbell  41.91
## accel_dumbbell_y 39.94
## magnet_belt_z  39.88
## accel_dumbbell_z 36.20
## magnet_belt_y  35.33
## roll_arm       33.70
## gyros_belt_z   31.27
## accel_forearm_x 28.88
## magnet_arm_x   28.58
## total_accel_dumbbell 28.42
## accel_arm_x    26.77
```

```
print(modFit.rf) # optimal mtry
```

```
## Random Forest
##
## 13737 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 13737, 13737, 13737, 13737, 13737, 13737, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
```

```
##      2      0.9877342  0.9844821
##     27      0.9876307  0.9843518
##     52      0.9743812  0.9675907
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
predict(modFit.rf,validate) # predict on validation data
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

We already fixed the tuning parameter `ntree=100`, but we can try another search for the optimal tuning parameter `mtry`. Below we use a different resampling method - 10 fold cross-validation repeated 3 times. This setup runs much longer, because of the three repeats, and provides a marginal improvement to the accuracy. The expected error is reduced to about 0.46%. The most important variables are the same, however the optimal `mtry` tuning parameter changes to more than 2 variables.

```
control <- trainControl(method="repeatedcv", number=10, repeats=3, search="random")
mtry <- sqrt(ncol(train))
```

```
pred.rf_random <- predict(modFit.rf_random,test) # predict on test data
confusionMatrix(pred.rf_random,test$classe) # measure accuracy on test data
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction    A    B    C    D    E
##      A 1673     4     0     0     0
##      B   1 1135     4     0     0
##      C    0     0 1020     9     3
##      D    0     0   2  955     5
##      E    0     0   0   0 1074
```

```
##
```

```
## Overall Statistics
```

```
##
##              Accuracy : 0.9952
##              95% CI : (0.9931, 0.9968)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##              Kappa : 0.994
```

```
## McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##              Class: A Class: B Class: C Class: D Class: E
## Sensitivity          0.9994   0.9965   0.9942   0.9907   0.9926
## Specificity          0.9991   0.9989   0.9975   0.9986   1.0000
## Pos Pred Value       0.9976   0.9956   0.9884   0.9927   1.0000
## Neg Pred Value       0.9998   0.9992   0.9988   0.9982   0.9983
## Prevalence           0.2845   0.1935   0.1743   0.1638   0.1839
```

```
## Detection Rate      0.2843  0.1929  0.1733  0.1623  0.1825
## Detection Prevalence 0.2850  0.1937  0.1754  0.1635  0.1825
## Balanced Accuracy   0.9992  0.9977  0.9958  0.9946  0.9963
```

```
varImp(modFit.rf_random) # variable importance
```

```
## rf variable importance
##
## only 20 most important variables shown (out of 52)
##
## Overall
## roll_belt      100.00
## yaw_belt       67.45
## pitch_forearm  63.01
## magnet_dumbbell_z 53.04
## magnet_dumbbell_y 51.49
## pitch_belt     49.78
## roll_forearm   45.20
## roll_dumbbell  27.98
## magnet_dumbbell_x 26.68
## accel_dumbbell_y 25.26
## magnet_belt_z  23.56
## accel_belt_z   23.25
## magnet_belt_y  21.37
## accel_dumbbell_z 19.29
## accel_forearm_x 19.26
## magnet_forearm_z 18.31
## roll_arm       16.00
## gyros_belt_z   15.16
## total_accel_dumbbell 15.05
## yaw_dumbbell   13.28
```

```
print(modFit.rf_random) # optimal mtry
```

```
## Random Forest
##
## 13737 samples
## 52 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 12364, 12362, 12364, 12363, 12361, 12363, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 8 0.9929389 0.9910675
## 9 0.9929631 0.9910979
## 12 0.9930359 0.9911903
## 25 0.9910460 0.9886728
## 27 0.9905121 0.9879977
## 28 0.9905606 0.9880589
## 29 0.9904391 0.9879053
```

```
## 38 0.9881584 0.9850198
## 41 0.9873817 0.9840372
## 42 0.9873812 0.9840359
## 44 0.9862410 0.9825941
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 12.
```

```
predict(modFit.rf_random,validate) # predict on validation data
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

## Appendix

The results from the lower performing models for this data set are presented below.

- Decision tree:

```
pred.tree <- predict(modFit.tree,test); confusionMatrix(pred.tree,test$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1524  463  488  440  147
##           B   27  376   33  169  156
##           C  119  300  505  355  302
##           D    0    0    0    0    0
##           E    4    0    0    0  477
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.4897
##           95% CI : (0.4769, 0.5026)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
##           Kappa : 0.3331
```

```
## McNemar's Test P-Value : NA
```

```
## Statistics by Class:
```

```
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9104 0.33011 0.49220 0.0000 0.44085
## Specificity      0.6348 0.91888 0.77856 1.0000 0.99917
## Pos Pred Value   0.4977 0.49409 0.31942      NaN 0.99168
## Neg Pred Value   0.9469 0.85109 0.87895 0.8362 0.88805
## Prevalence       0.2845 0.19354 0.17434 0.1638 0.18386
## Detection Rate   0.2590 0.06389 0.08581 0.0000 0.08105
## Detection Prevalence 0.5203 0.12931 0.26865 0.0000 0.08173
## Balanced Accuracy 0.7726 0.62450 0.63538 0.5000 0.72001
```

- Generalized boosted regression (gbm):

```
pred.gbm <- predict(modFit.gbm,test); confusionMatrix(pred.gbm,test$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1647   37    0    1    2
##           B   16 1067   28    1   19
##           C    7   32  980   27    9
##           D    4    3   16  933   24
##           E    0    0    2    2 1028
##
## Overall Statistics
##
##           Accuracy : 0.9609
##           95% CI : (0.9556, 0.9657)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9506
##           McNemar's Test P-Value : 3.597e-10
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9839   0.9368   0.9552   0.9678   0.9501
## Specificity           0.9905   0.9865   0.9846   0.9904   0.9992
## Pos Pred Value        0.9763   0.9434   0.9289   0.9520   0.9961
## Neg Pred Value        0.9936   0.9849   0.9905   0.9937   0.9889
## Prevalence            0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate        0.2799   0.1813   0.1665   0.1585   0.1747
## Detection Prevalence  0.2867   0.1922   0.1793   0.1665   0.1754
## Balanced Accuracy      0.9872   0.9617   0.9699   0.9791   0.9746
```

- Linear discrimination analysis:

```
pred.lda <- predict(modFit.lda,test); confusionMatrix(pred.lda,test$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1386  173  106  42   29
##           B   41  728   99  34  200
##           C  118  136  664 112   80
##           D  125   52  128 733  121
##           E    4   50   29  43  652
##
## Overall Statistics
##
```

```

##                Accuracy : 0.7074
##                95% CI : (0.6956, 0.719)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##                Kappa : 0.6298
##  McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##                Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8280   0.6392   0.6472   0.7604   0.6026
## Specificity      0.9169   0.9212   0.9082   0.9134   0.9738
## Pos Pred Value   0.7984   0.6606   0.5982   0.6324   0.8380
## Neg Pred Value   0.9306   0.9141   0.9242   0.9511   0.9158
## Prevalence       0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate   0.2355   0.1237   0.1128   0.1246   0.1108
## Detection Prevalence 0.2950   0.1873   0.1886   0.1969   0.1322
## Balanced Accuracy 0.8724   0.7802   0.7777   0.8369   0.7882

```